

ENSF 612

Lecture – PySpark Programming

Dr. Gias Uddin, Assistant Professor
Electrical and Software Engineering
University of Calgary

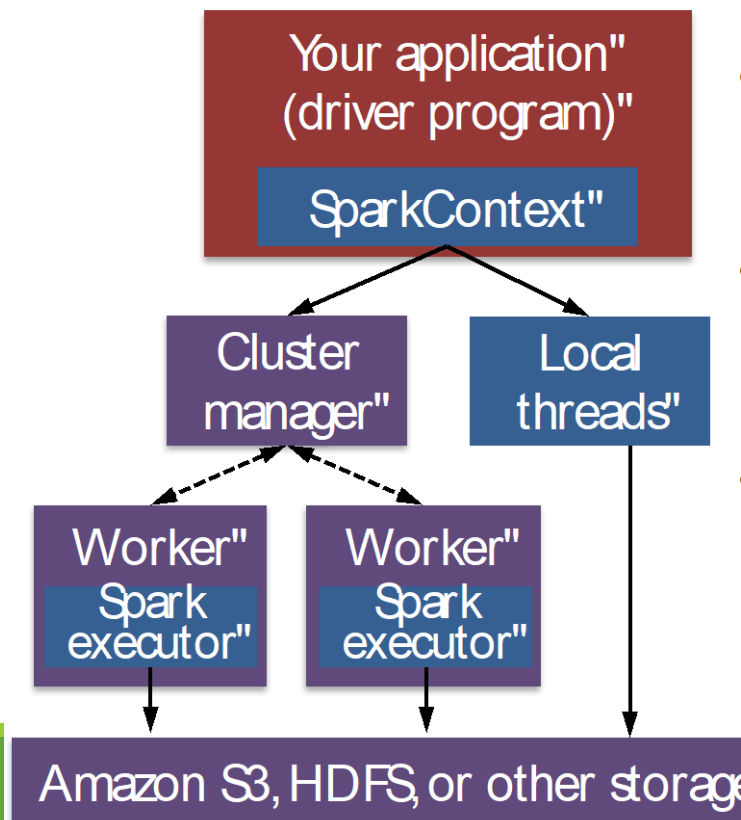
<https://giasuddin.ca/>

Spark programming

- ◆ Will use the Python interface to Spark pySpark
- ◆ PySpark code let us do the following:
 - ⑩ “Here are some operations. Apply them to my data”
- ◆ RDDs are the key concept
 - RDD – Resilient Distributed Dataset
 - RDD is a collection of data elements partitioned across the nodes of a cluster that can be operated in parallel

Spark programming – cont'd

- ◆ Spark program is two programs
 - ⑩ A **driver** program and a **workers** program
- ◆ Worker programs run on cluster nodes typically
- ◆ RDDs are distributed across workers



Spark programming – cont'd

◆ Initialization of Spark

```
import org.apache.spark.SparkContext  
import org.apache.spark.SparkConf
```

```
conf = SparkConf().setAppName(appName).setMaster(master)  
sc = SparkContext(conf=conf)
```

◆ A Spark program first creates a ***SparkContext*** object

⑩ Tells Spark how and where to access a cluster

⑩ pySpark shell automatically creates ***sc*** variable

⑩ iPython programs must use a constructor to create ***SparkContext***

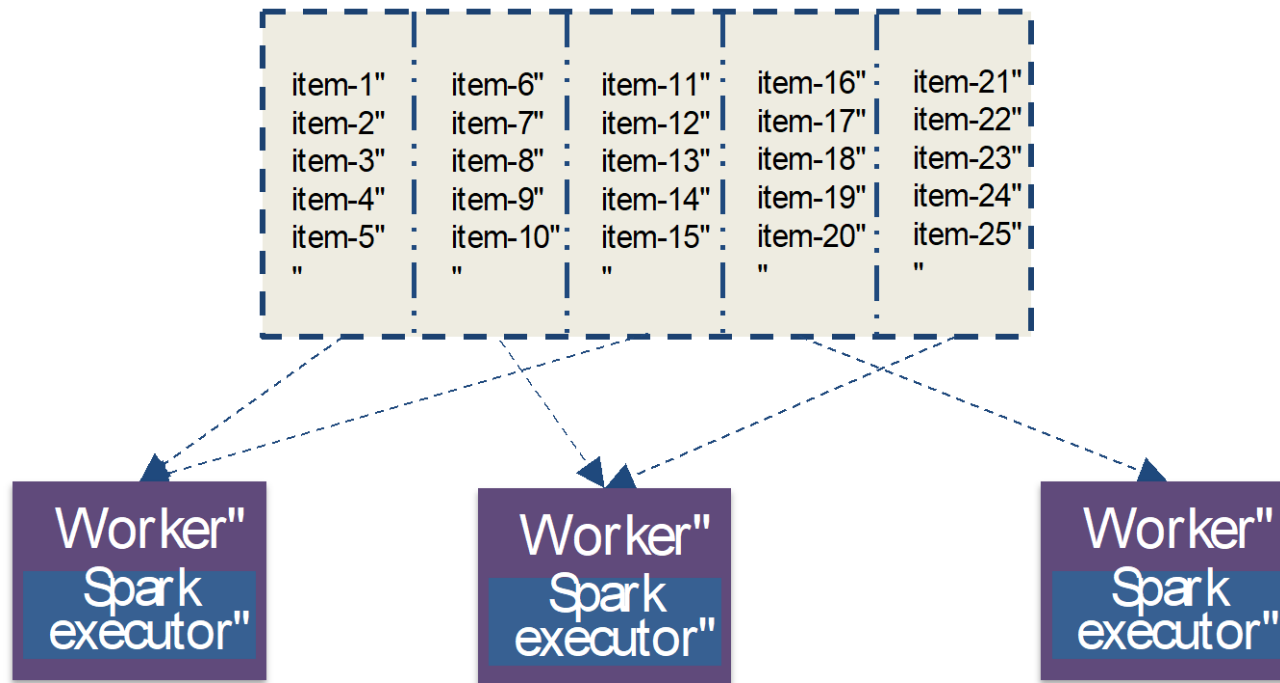
◆ ***SparkContext*** object can be used to create RDDs

Spark programming – cont'd

- ◆ RDDs are immutable once they are constructed
 - ⑩ Can transform them to other RDDs
 - ⑩ Can perform actions on them, e.g., read elements
 - ⑩ However, can't change them
- ◆ RDDs allow parallel operations on collections
- ◆ How do we construct RDDs?
 - ◆ From Python collections (lists)
 - ◆ By transforming other RDDs
 - ◆ From files stored in HDFS and other storage

Spark programming – cont'd

- ◆ How is an RDD parallelized?
- ◆ RDD can be split into programmer specified # of partitions
- ◆ More partitions – more parallelism



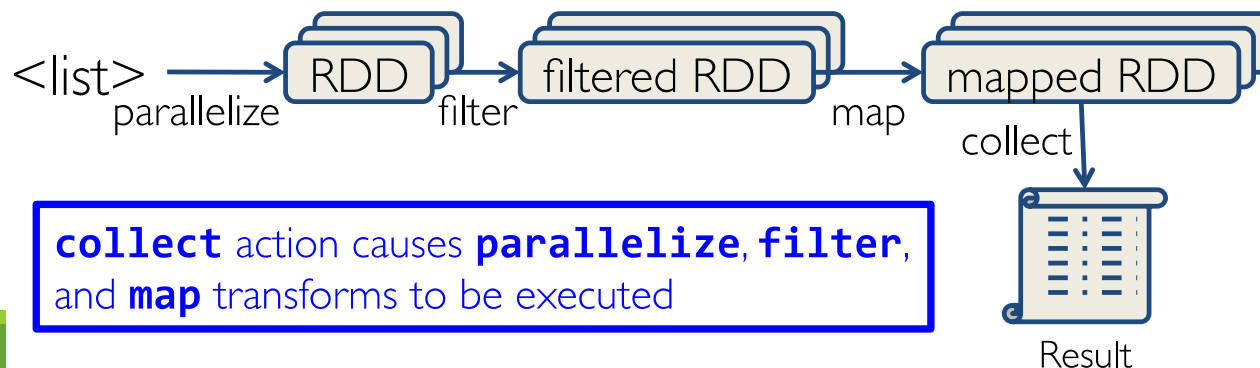
Spark programming – cont'd

- ◆ Two types of operations – *transformations* and *actions*
- ◆ Transformations are lazily evaluated for efficiency
 - ⑩ Transformations execute only when action performed on RDD
- ◆ RDDs can be persisted (cached) in memory or disk

Create an RDD from a data source:  <list>

Apply transformations to an RDD: map filter

Apply actions to an RDD: collect count



Spark programming – cont'd

- ◆ Will now look at transformations and actions in Spark
- ◆ Need to review Python **lambda** functions first
- ◆ Small anonymous functions not bound to a name

lambda a,b: a+b

Returns sum of two argument a and b

- ◆ Restricted to single expressions

Spark programming – cont'd

◆ Frequently used Spark *transformations*

Transformation	Description
<code>map(func)</code>	return a new distributed dataset formed by passing each element of the source through a function <i>func</i>
<code>filter(func)</code>	return a new dataset formed by selecting those elements of the source on which <i>func</i> returns true
<code>distinct([numTasks])</code>	return a new dataset that contains the distinct elements of the source dataset
<code>flatMap(func)</code>	similar to map, but each input item can be mapped to 0 or more output items (so <i>func</i> should return a Seq rather than a single item)

Spark programming – cont'd

◆ Example *transformations*

```
>>> rdd = sc.parallelize([1, 2, 3, 4])
>>> rdd.map(lambda x: x * 2)
RDD: [1, 2, 3, 4] → [2, 4, 6, 8]
```

Function literals (green)
are closures automatically
passed to workers

```
>>> rdd.filter(lambda x: x % 2 == 0)
RDD: [1, 2, 3, 4] → [2, 4]
```

```
>>> rdd2 = sc.parallelize([1, 4, 2, 2, 3])
>>> rdd2.distinct()
RDD: [1, 4, 2, 2, 3] → [1, 4, 2, 3]
```

```
>>> rdd = sc.parallelize([1, 2, 3])
>>> rdd.Map(lambda x: [x, x+5])
RDD: [1, 2, 3] → [[1, 6], [2, 7], [3, 8]]
```

```
>>> rdd.flatMap(lambda x: [x, x+5])
RDD: [1, 2, 3] → [1, 6, 2, 7, 3, 8]
```

Spark programming – cont'd

- ◆ Spark *action* causes transformations to be applied
- ◆ Mechanism for getting results out of Spark
- ◆ Frequently used *actions* in Spark

Action	Description
<code>reduce(func)</code>	aggregate dataset's elements using function <i>func</i> . <i>func</i> takes two arguments and returns one, and is commutative and associative so that it can be computed correctly in parallel
<code>take(n)</code>	return an array with the first <i>n</i> elements
<code>collect()</code>	return all the elements as an array WARNING: make sure will fit in driver program
<code>takeOrdered(n, key=func)</code>	return <i>n</i> elements ordered in ascending order or as specified by the optional key function

Spark programming – cont'd

◆ Examples of *actions*

```
>>> rdd = sc.parallelize([1, 2, 3])  
>>> rdd.reduce(lambda a, b: a * b)  
Value: 6
```

```
>>> rdd.take(2)  
Value: [1,2] # as list
```

```
>>> rdd.collect()  
Value: [1,2,3] # as list
```

```
>>> rdd = sc.parallelize([5,3,1,2])  
>>> rdd.takeOrdered(3, lambda s: -1 * s)  
Value: [5,3,2] # as list
```

Spark programming – cont'd

◆ Lets put transformations and actions together

#Create RDD that has 4 partitions

#Each element of RDD is a line of the text file

lines = sc.textFile("...",4)

#isNotComment function only outputs a line if it is not a #comment

noComments = lines.filter(isNotComment)

print lines.count()



count() causes Spark to:

- read data
- sum within partitions
- combine sums in driver

Spark programming – cont'd

- ◆ Spark may be forced to read data multiple times

#Create RDD that has 4 partitions

#Each element of RDD is a line of the text file

lines = sc.textFile(...,4)

#isNotComment function only outputs a line if it is not a #comment

noComments = lines.filter(isNotComment)

print lines.count(), noComments.count()

Spark recomputes lines:

- read data (again)
- sum within partitions
- combine sums in driver

Spark programming – cont'd

- ◆ Spark can persist (cache) RDDs to be more efficient

#Create RDD that has 4 partitions

#Each element of RDD is a line of the text file

lines = sc.textFile(...,4)

#save lines in memory – don't recompute!

lines.cache()

#isNotComment function only outputs a line if it is not a #comment

noComments = lines.filter(isNotComment)

print lines.count(), noComments.count()

Spark programming – cont'd

- ◆ Lifecycle of a Spark program
- ◆ Create RDDs from external data or Python collections
- ◆ Lazily transform them into new RDDs
- ◆ Cache some RDDs for reuse
- ◆ Perform actions to trigger parallel computations and results

Spark programming – cont'd

- ◆ Similar to MapReduce, Spark supports Key-Value pairs
- ◆ Spark allows **pair RDDs** to be created
- ◆ Each element is a tuple consisting of a key and value

```
>>> rdd = sc.parallelize([(1, 2), (3, 4)])  
RDD: [(1, 2), (3, 4)]
```

Spark programming – cont'd

◆ Examples of Key-Value transformations

Key-Value Transformation	Description
reduceByKey(<i>func</i>)	return a new distributed dataset of (K,V) pairs where the values for each key are aggregated using the given reduce function <i>func</i> , which must be of type (V,V) → V
sortByKey()	return a new dataset (K,V) pairs sorted by keys in ascending order
groupByKey()	return a new dataset of (K, Iterable<V>) pairs

Spark programming – cont'd

◆ Examples of Key-Value transformations

```
>>> rdd = sc.parallelize([(1,2), (3,4), (3,6)])
```

```
>>> rdd.reduceByKey(lambda a, b: a + b)
```

```
RDD: [(1,2), (3,4), (3,6)] → [(1,2), (3,10)]
```

```
>>> rdd2 = sc.parallelize([(1,'a'), (2,'c'), (1,'b')])
```

```
>>> rdd2.sortByKey()
```

```
RDD: [(1,'a'), (2,'c'), (1,'b')] →  
      [(1,'a'), (1,'b'), (2,'c')]
```

```
>>> rdd2.groupByKey()
```

```
RDD: [(1,'a'), (1,'b'), (2,'c')] →  
      [(1,['a','b']), (2,['c'])]
```

Be careful using `groupByKey()` – can cause a lot of data movement across the network!

Spark programming – cont'd

- ◆ Spark supports working with structured data

- ◆ Supports *join* operations on pair RDDs

- ◆ **X.join(Y)**

- ◆ Return RDD of all pairs of elements with matching keys in X and Y

- ◆ Each pair is (k,(v1,v2)) tuple, where (k, V1) is in X and (k,V2) is in Y

```
>>> x = sc.parallelize([("a", 1), ("b", 4)])  
>>> y = sc.parallelize([("a", 2), ("a", 3)])  
>>> sorted(x.join(y).collect())
```

```
Value: [('a', (1, 2)), ('a', (1, 3))]
```

Spark programming – cont'd

◆ **X.leftOuterJoin(Y)**

- ◆ For each element (k,v) in X, resulting RDD will either contain
 - ◆ All pairs of (k,(v,w)) for w in Y
 - ◆ Or the pair (k,(v,None)) if no element in Y has key k

```
>>> x = sc.parallelize([("a", 1), ("b", 4)])  
>>> y = sc.parallelize([("a", 2)])  
>>> sorted(x.leftOuterJoin(y).collect())
```

```
Value: [('a', (1, 2)), ('b', (4, None))]
```

Spark programming – cont'd

◆ **Y.rightOuterJoin(X)**

- ◆ For each element (k,w) in X , resulting RDD will either contain
 - ◆ All pairs of $(k,(v,w))$ for v in Y
 - ◆ Or the pair $(k,(None,w))$ if no element in Y has key k

```
>>> x = sc.parallelize([("a", 1), ("b", 4)])  
>>> y = sc.parallelize([("a", 2)])  
>>> sorted(y.rightOuterJoin(x).collect())
```

```
Value: [('a', (2, 1)), ('b', (None, 4))]
```

Spark programming – cont'd

◆ **Y.fullOuterJoin(X)**

- ◆ For each element (k,v) in X, resulting RDD will either contain
 - ◆ All pairs (k,(v,w)) for w in Y, or (k,(v,None)) if no elements in Y have k
- ◆ For each element (k,w) in Y, resulting RDD will either contain
 - ◆ All pairs (k,(v,w)) for v in X, or (k,(None,w)) if no elements in X have k

```
>>> x = sc.parallelize([("a", 1), ("b", 4)])  
>>> y = sc.parallelize([("a", 2), ("c", 8)])  
>>> sorted(x.fullOuterJoin(y).collect())
```

```
Value: [('a', (1, 2)), ('b', (4, None)) , ('c', (None, 8))]
```

Spark programming – cont'd

- ◆ Spark **shared variables**
- ◆ Spark automatically creates closures for:
 - ⑩ Functions that run on RDDs at workers
 - ⑩ Global variables used by those workers
- ◆ One closure per worker
 - ⑩ Sent for every task
 - ⑩ No communication between workers
 - ⑩ Changes to global variables at workers not sent to drivers
- ◆ **This model may be inefficient for many use cases**

Spark programming – cont'd

- ◆ Consider following 2 use cases
 - ⑩ Sending large lookup table to workers
 - ⑩ Sending large feature vectors in a ML algorithm to workers
- ◆ Counting events that occur during job execution
 - ⑩ E.g., How many input lines were blank?
 - ⑩ E.g., How many input records were corrupt?
- ◆ **Problems**
 - ◆ Closures are (re)sent with every job
 - ◆ Inefficient to send large data to each worker
 - ◆ Closures are one way – driver → worker

Spark programming – cont'd

- ◆ **Broadcast** and **accumulator** variables address this

- ◆ Broadcast variables

 - ⑩ Efficiently send large, **read-only** values to all workers

 - ⑩ Saved at workers for use in one or more Spark operations

- ◆ Accumulators

 - ◆ Aggregate values from workers back to the driver

 - ◆ Only driver can access the value of accumulator

 - ◆ For tasks, accumulators are **write-only**

Spark programming – cont'd

◆ Broadcast variables

- ◆ Ship to each worker only once instead of with each task
- ◆ Efficiently give each worker a large dataset
- ◆ Distributed using efficient broadcast algorithms

At the driver:

```
>>> broadcastVar = sc.broadcast([1, 2, 3])
```

At a worker (in code passed via a closure)

```
>>> broadcastVar.value4  
[1, 2, 3]
```

Spark programming – cont'd

◆ Broadcast variables example

Country code lookup for HAM radio call signs"

```
#, Lookup, the, locations, of, the, call, signs, on, the,  
#, RDD, contactCounts. , We, load, a, list, of, call, sign, ,  
#, prefixes, to, country, code, to, support, this, lookup, ,  
signPrefixes = loadCallSignTable()
```

Expensive to send large table"
(Re-)sent for every processed file "

```
def4 processSignCount(sign_count, signPrefixes):  
    country = lookupCountry(sign_count[0], signPrefixes)  
    count = sign_count[1]  
    return4(country, count)
```

```
countryContactCounts = (contactCounts  
    .map(processSignCount)  
    .reduceByKey((lambda4x, y: x+ y)))
```

From: <http://shop.oreilly.com/product/0636920028512.do> "

Spark programming – cont'd

◆ Broadcast variables example

```
#, Lookup, the, locations, of, the, call, signs, on, the,  
#, RDD, contactCounts. , We, load, a, list, of, call, sign, ,  
#, prefixes, to, country, code, to, support, this, lookup, ,  
signPrefixes = sc.broadcast(loadCallSignTable())
```

Efficiently sent once to workers"

```
def4 processSignCount(sign_count, signPrefixes):  
    country = lookupCountry(sign_count[0], signPrefixes.value)  
    count = sign_count[1]  
    return4(country, count)
```

```
countryContactCounts = (contactCounts  
    .map(processSignCount)  
    .reduceByKey((lambda4x, y: x+ y)))
```

From: <http://shop.oreilly.com/product/0636920028512.do> "

Spark programming – cont'd

◆ Accumulators

- ◆ Variables that can only be “added” to by associative operation
- ◆ Used to efficiently implement parallel counters and sum
- ◆ Only driver can read accumulator value – not tasks

```
>>> accum = sc.accumulator(0)
>>> rdd = sc.parallelize([1, 2, 3, 4])
>>> def f(x):
>>>     global accum
>>>     accum += x

>>> rdd.foreach(f)
>>> accum.value
Value: 10
```

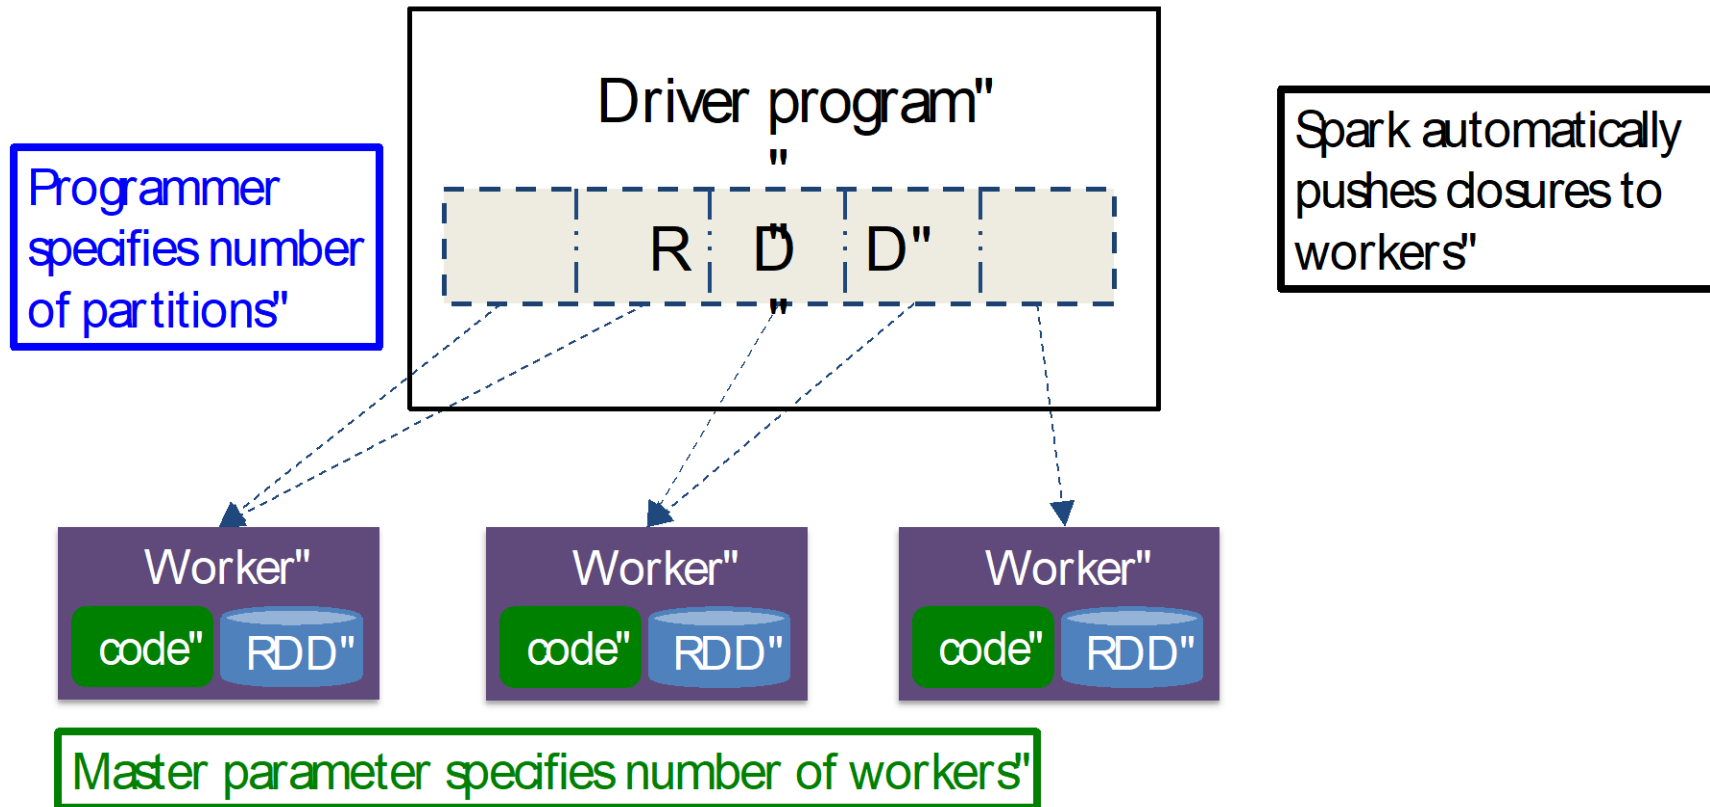
Spark programming – cont'd

◆ Accumulators

- ◆ Tasks at workers can't access accumulator's values
- ◆ Tasks see accumulators as write-only variables
- ◆ Can be used in actions or transformations
 - ◆ Actions – each update to accumulator applied only once
 - ◆ Transformations – no such guarantees! (why?)

Spark programming – cont'd

◆ Summary of Spark programming model



Acknowledgements

Portions of these slides were adapted from external material available under creative commons license CC-BY-NC-SA 4.0. This license grants the ability to share and adapt the material for non-commercial purposes.

Name of the creator: Dr. Anthony Joseph, University of Berkeley and team

License notice: <https://creativecommons.org/licenses/by-nc-sa/4.0/legalcode>

Copyright notice: CC-BY-NC-SA 4.0

Link to material: <https://courses.edx.org/courses/BerkeleyX/CS100.1x/1T2015/info>