

ENSF 612

Lecture MapReduce Algorithms

Dr. Gias Uddin,
Electrical and Software Engineering,
University of Calgary.
<https://giasuddin.ca/>

Topics

- ◆ Relational operations
- ◆ Matrix operations
- ◆ Graph algorithms

Relational operations

- ◆ Can do SQL-like operations using MapReduce
- ◆ Consider *select* and *join* operations as examples
- ◆ Recap some terminology first
- ◆ **Relation (R)** – a table with names of its columns
- ◆ **Attributes** – names of columns
- ◆ **Tuples** – rows of a relation
- ◆ **Schema** – collection of attributes
 - ◆ denoted by $R(A_1, A_2, \dots, A_N)$ where the A_i s are attributes

Relational operators – cont'd

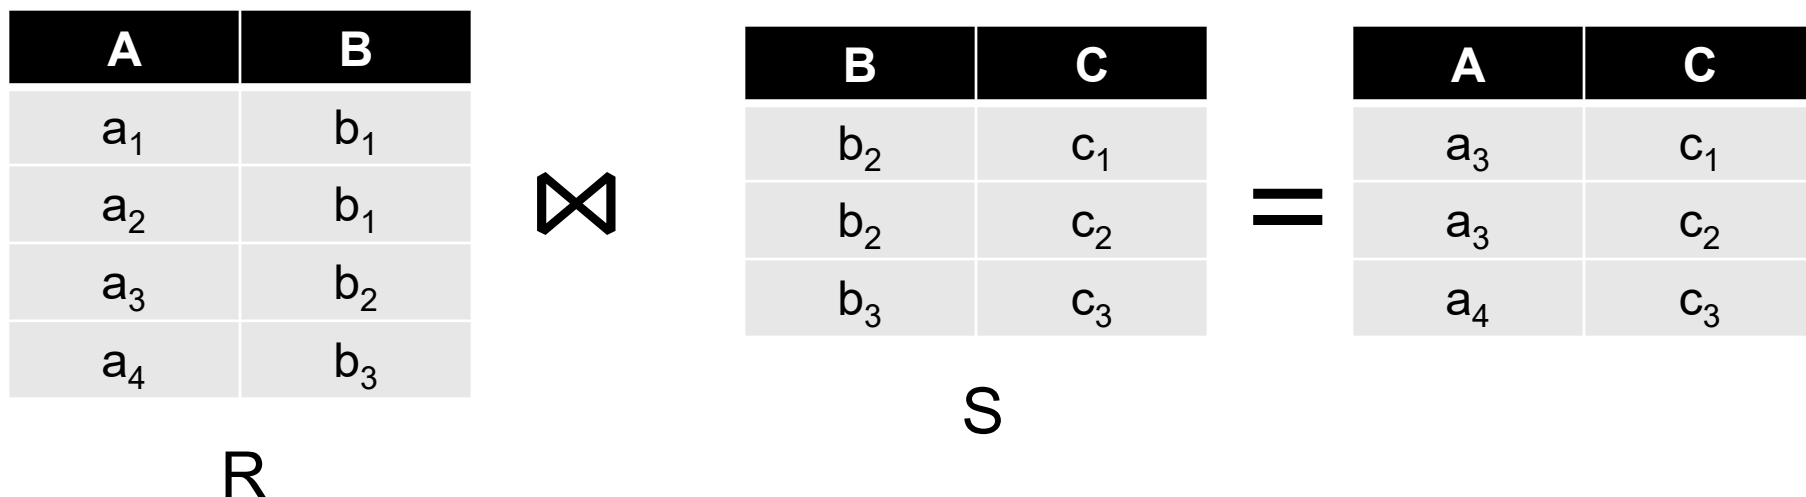
- ◆ How can we implement *select* using MapReduce?

*select * from EMPLOYEES where salary > 50,000*

- ◆ Apply condition C on each tuple in a relation
- ◆ Can implement using only a map step
- ◆ Input for **Map** is $\langle t, t \rangle$ where t is a tuple in the relation
- ◆ **Map**($\langle t, t \rangle$) outputs $\langle t, t \rangle$ if C is satisfied; nothing if not.
- ◆ **Reduce** is an identity function – simply emits input from **Map**
- ◆ Can either use key or value from **Reduce** as result

Relational operators– cont'd

- ◆ How to implement a *join* using MapReduce?
- ◆ **Compute the natural join** $R(A,B) \bowtie S(B,C)$
 - ◆ R and S are each stored in files
 - ◆ Tuples are pairs (a,b) or (b,c)



Relational operators– cont'd

- ◆ **Use a hash function h from B-values to $1...k$**
- ◆ **A Map process turns:**
 - ◆ Each input tuple $R(a,b)$ into key-value pair $(b,(a,R))$
 - ◆ Each input tuple $S(b,c)$ into $(b,(c,S))$
- ◆ **Map processes**
 - ◆ send each k-v pair with key b to Reduce process $h(b)$
 - ◆ Hadoop does this automatically; just tell it what k is.
- ◆ **Each Reduce process**
 - ◆ matches all the pairs $(b,(a,R))$ with all $(b,(c,S))$
 - ◆ outputs (a,b,c) .

Graph algorithms

- ◆ Graph algorithms have many applications
 - ◆ Social network analysis
 - ◆ Shortest path in road networks
- ◆ Consider the breadth first search (BFS) algorithm
- ◆ Traverse a graph in a breadth first manner
- ◆ Many applications – e.g., calculating shortest path
- ◆ How can we implement BFS using MapReduce?

Graph algorithms – cont'd

- ◆ BFS terminology
- ◆ **Graph** consists of **nodes**
- ◆ Two nodes may be connected by an **edge**
- ◆ Designate one node as *source*
- ◆ Nodes have **colours**
 - ◆ **White (W)** – node has not been visited
 - ◆ **Grey (G)** – node visited but its kids not visited, i.e., W
 - ◆ **Black (B)** – node visited all its kids are G
- ◆ **Goal – start from source -stop when no G nodes**

BFS

Wikipedia:

https://en.wikipedia.org/wiki/Breadth-first_search

Breadth-first search (BFS) is an [algorithm](#) for traversing or searching [tree](#) or [graph](#) data structures. It starts at the [tree root](#) (or some arbitrary node of a graph, sometimes referred to as a 'search key'[\[1\]](#)), and explores all of the neighbor nodes at the present depth prior to moving on to the nodes at the next depth level.

It uses the opposite strategy of [depth-first search](#), which instead explores the node branch as far as possible before being forced to backtrack and expand other nodes.[\[2\]](#)

Pseudocode [\[edit\]](#)


Input: A graph G and a starting vertex $root$ of G

Output: Goal state. The *parent* links trace the shortest path back to $root$ [\[7\]](#)

```
1  procedure BFS( $G$ ,  $root$ ) is
2      let  $Q$  be a queue
3      label  $root$  as discovered
4       $Q.enqueue(root)$ 
5      while  $Q$  is not empty do
6           $v := Q.dequeue()$ 
7          if  $v$  is the goal then
8              return  $v$ 
9          for all edges from  $v$  to  $w$  in
 $G.adjacentEdges(v)$  do
10             if  $w$  is not labeled as discovered
then
11                 label  $w$  as discovered
12                  $Q.enqueue(w)$ 
```

Queue – First in First Out (FIFO)
Stack – Last in First Out (LIFO)

BFS




BFS


- ◆ Using BFS to get shortest path from *source* to nodes
- ◆ Need a couple of extra variables for a node
- ◆ **Score** – in how many hops can I get to *source* node
 - ◆ Initially 0 for *source* and infinity for others
- ◆ **Parent** – from which parent did I visit this node
 - ◆ Initially null for all nodes
- ◆ How does the algorithm work?

- ◆ Start from *source* (process all "*G*" nodes)
 - ◆ Mark from "*G*" to "*B*"
 - ◆ Visit unvisited (i.e., "*W*") kids - mark such kids from "*W*" to "*G*"
 - ◆ Set *parent* of kids as the node from which visited, i.e., *source*
 - ◆ Set *score* of kids as score of *parent* + 1, i.e., 0+1
- ◆ Repeat process for all the new "*G*" nodes
- ◆ Keep going till all nodes reachable from *source* are "*B*"
- ◆ We visit nodes with *score* $K+1$ from those with *score* K
- ◆ Score of destination gives shortest hop to source
- ◆ Can also trace shortest path by using *parent* variable


BFS




BFS



BFS



BFS



Finding the Shortest Path using BFS in MapReduce

Intuition

1. $\text{DistanceTo}(\text{startNode}) = 0$
2. For all nodes n directly reachable from startNode , $\text{DistanceTo}(n) = 1$
3. For all nodes n reachable from some other set of nodes S ,
$$\text{DistanceTo}(n) = 1 + \min(\text{DistanceTo}(m), m \in S)$$

Algorithm

1. A map task receives a node n as a key, and $(D, \text{points-to})$ as its value
 - D is the distance to the node from the start
 - “points-to” is a list of nodes reachable from n
 - $\forall p \in \text{points} - \text{to}, \text{emit}(p, D + 1)$
2. Reduce task gathers possible distances to a given p and selects the minimum one

BFS using MapReduce

- ◆ How to implement parallel BFS with MapReduce?
- ◆ Assume we have large input file describing graph
- ◆ Format of input file is as follows
- ◆ *node<tab>adjacency_list/ score/ colour/ parentNode*
- ◆ *node* is the key, i.e., node id - everything else is a value
- ◆ *adjacency_list* – comma separated list of kid nodes
- ◆ *parentNode* – “*source*” for *source* node, *null* for others
- ◆ Algorithm is iterative
 - ◆ Do Map-Reduce pair – produce a result
 - ◆ Result is input to another Map-Reduce pair
 - ◆ Stop when all nodes processed

BFS using MapReduce

◆ Map step

- ◆ Node is " W " – just emit the line
- ◆ Node is " B " – just emit the line
- ◆ Node is " G "
 - ◆ Make node " B " – emit a line corresponding to this
 - ◆ Make kids of node " G "
 - ◆ For each kid node, set parent as the node
 - ◆ For each kid node, set *score* as parent node *score* + 1
 - ◆ Emit records corresponding to each kid node
 - ◆ How do we handle kids' adjacency lists?

BFS using MapReduce

◆ Reduce step

- ◆ Aggregate records with same key, i.e., node id
- ◆ Will get multiple records for same node
- ◆ Need to aggregate them and emit one valid record
- ◆ A valid record should contain
 - ◆ a valid adjacency list for the node
 - ◆ the most up to date colour, i.e., state, of the node
 - ◆ the correct score and parent

BFS using MapReduce

◆ Reduce step

◆ Questions to ponder

- ◆ Why will there be multiple records for same node?
- ◆ Why will the adjacency list of some records be empty?
- ◆ Why will we get a node marked as both " W " and " G "
- ◆ Why will we get a node marked as both " G " and " B "?
- ◆ Why will we get 2 different scores and parents for same node?

◆ Reducer is somewhat more complex than the mapper

◆ There are libraries which simply do graph processing

◆ Pregel, graphX

- ◆ Relational operations
- ◆ Matrix operations
- ◆ Graph algorithms

Acknowledgements

Portions of these slides were adapted with permission from Leskovec et al.

(<http://www.mmds.org>)

Portions of these slides were adapted with permission as per the following details:

License notice: <http://creativecommons.org/licenses/by-sa/3.0/>

Copyright notice: CC-BY-SA 3.0

Link to material:

<https://hadooptutorial.wikispaces.com/Iterative+MapReduce+and+Counters>